

Jack Devlin

Platform & AI Infrastructure Engineer · Northern Ireland
Irish & British citizen, EU and UK work rights · open to relocation
07423 281519 · jack@devlinops.com · devlinops.com

PROFILE

Platform engineer specialising in the infrastructure AI products run on: Kubernetes, self-hosted observability and CI/CD underneath, LLM gateways and guardrails on top. Three years on a high-volume microservices estate, most recently taking natural-language AI tools into production. Completing an MSc in AI: a capacity-aware recovery scheduler for Kubernetes, measured on real EKS clusters under induced node failure. Long-form technical case studies at devlinops.com.

EXPERIENCE

Loweconex (via Devlinops Ltd)

Lisburn, UK · Aug 2023 – Present

Platform & Site Reliability Engineer

Core platform engineering, on-call rotation and AI infrastructure for a high-volume microservices environment. Progressed from graduate QA to platform engineering; moved to independent B2B contract in September 2025.

- **Clarity AI.** Production infrastructure and trust layer for a natural-language database query tool on Google Gemini, live across ~30 tenants for ~20 daily users. Verbatim SQL on every answer, five classes of hallucination caught mechanically each turn, nightly-compiled schema intelligence injected into prompts, circuit breakers for runaway agent loops, and canary evals gating every deploy.
- **AI platform.** Own the internal LLM gateway: every AI workload behind one endpoint, virtual keys with model allowlists that fail closed, spend attributed to the tenant and feature that caused it. Built an automated pull request review agent on top of it, running on PRs across eight production services.
- **Pipeline platform.** A Bitbucket shared pipeline library for Java and Node, replacing bespoke per-service pipelines. One `.ci/builds.yaml` per service; Veracode SAST and SourceClear SCA scans file findings straight to Jira. Sustains ~400 deployments a month across 20 services, verified after deploy by Sentry, its companion test framework.
- **Self-hosted observability.** Metrics through Prometheus and Thanos, logs in Loki, traces in Tempo, alerting via Alertmanager. Set the OpenTelemetry tracing and trace-to-log correlation patterns developers build against. 20+ dashboards as code and 50+ runbook-backed alerts over 20 services and four environments.
- **Heimdall.** A deployment intelligence dashboard answering “where is my ticket right now?”. Aggregates ArgoCD, Bitbucket, Kubernetes, Jira and test runs to surface environment drift, PR bottlenecks and DORA metrics. TimescaleDB backed; pages read a pre-computed in-memory snapshot, never the database. Used daily by 20+ engineers.

OD3 Engineering

Magherafelt, UK · Aug 2021 – Aug 2022

Software Developer (Industrial Placement)

- Built Tekla Structures API applications and design-platform integrations, automating structural-design workflows to cut manual drafting time for in-house architects.

TECHNICAL

Platform & IaC: Kubernetes, EKS, ArgoCD (App-of-Apps, Image Updater), Helm, Kustomize, Bitbucket Pipelines, Terraform, AWS CDK.

AI & ML: LLM gateways (LiteLLM), Google Gemini API, LLM guardrails and evaluation, schema-grounded retrieval, PyTorch, Stable-Baselines3, PPO.

Observability & Security: Prometheus, Grafana, Loki, Thanos, Tempo, Alertmanager, OpenTelemetry, Falco, Suricata, Veracode.

Cloud, Data & Languages: AWS (EKS, Lambda, IAM, VPC, S3), PostgreSQL, TimescaleDB, Kafka, Redis; Python, Java, TypeScript, Bash.

EDUCATION

Queen's University Belfast

Belfast, UK

MSc Artificial Intelligence (Expected Sep 2026) · BSc Computer Science (2:1), first-class dissertation on deep learning for stock-price prediction (2019 – 2023)

- MSc dissertation: capacity-aware recovery scheduling for Kubernetes after node failure. Evaluated on real EKS clusters against kube-scheduler and PriorityClass baselines, with a pre-registered analysis plan; workloads shaped from the Alibaba 2018 trace.